

Personalizing Image Search Results on Flickr

Kristina Lerman, Anon Plangprasopchok and Chio Wong

University of Southern California
Information Sciences Institute
4676 Admiralty Way
Marina del Rey, California 90292
{lerman,plangpra,chiowong}@isi.edu

Abstract

The social media site Flickr allows users to upload their photos, annotate them with tags, submit them to groups, and also to form social networks by adding other users as contacts. Flickr offers multiple ways of browsing or searching it. One option is tag search, which returns all images tagged with a specific keyword. If the keyword is ambiguous, e.g., “beetle” could mean an insect or a car, tag search results will include many images that are not relevant to the sense the user had in mind when executing the query. We claim that users express their photography interests through the metadata they add in the form of contacts and image annotations. We show how to exploit this metadata to personalize search results for the user, thereby improving search performance. First, we show that we can significantly improve search precision by filtering tag search results by user’s contacts or a larger social network that includes those contact’s contacts. Secondly, we describe a probabilistic model that takes advantage of tag information to discover latent topics contained in the search results. The users’ interests can similarly be described by the tags they used for annotating their images. The latent topics found by the model are then used to personalize search results by finding images on topics that are of interest to the user.

Introduction

The photosharing site Flickr is one of the earliest and more popular examples of the new generation of Web sites, labeled *social media*, whose content is primarily user-driven. Other examples of social media include: blogs (personal online journals that allow users to share thoughts and receive feedback on them), Wikipedia (a collectively written and edited online encyclopedia), and Del.icio.us and Digg (Web sites that allow users to share, discuss, and rank Web pages, and news stories respectively). The rise of social media underscores a transformation of the Web as fundamental as its birth. Rather than simply searching for, and passively consuming, information, users are collaboratively creating, evaluating, and distributing information. In the near future, new information-processing applications enabled by social media will include tools for personalized information discovery, applications that exploit the “wisdom of crowds” (e.g., emergent semantics and collaborative in-

formation evaluation), deeper analysis of community structure to identify trends and experts, and many others still difficult to imagine.

Social media sites share four characteristics: (1) Users create or contribute content in a variety of media types; (2) Users annotate content with tags; (3) Users evaluate content, either actively by voting or passively by using content; and (4) Users create social networks by designating other users with similar interests as contacts or friends. In the process of using these sites, users are adding rich metadata in the form of social networks, annotations and ratings. Availability of large quantities of this metadata will lead to the development of new algorithms to solve a variety of information processing problems, from new recommendation to improved information discovery algorithms.

In this paper we show how user-added metadata on Flickr can be used to improve image search results. We claim that users express their photography interests on Flickr, among other ways, by adding photographers whose work they admire to their social network and through the tags they use to annotate their own images. We show how to exploit this information to personalize search results to the individual user.

The rest of the paper is organized as follows. First, we describe tagging and why it can be viewed as a useful expression of user’s interests, as well as some of the challenges that arise when working with tags. In Section “Anatomy of Flickr” we describe Flickr and its functionality in greater details, including its tag search capability. In Section “Data collections” we describe the data sets we have collected from Flickr, including image search results and user information. In Sections “Personalizing by contacts” and “Personalizing by tags” we present the two approaches to personalize search results for an individual user by filtering by contacts and filtering by tags respectively. We evaluate the performance of each method on our Flickr data sets. We conclude by discussing results and future work.

Tagging for organizing images

Tags are keyword-based metadata associated with some content. Tagging was introduced as a means for users to organize their own content in order to facilitate searching and browsing for relevant information. It was popularized by

the social bookmarking site Delicious¹, which allowed users to add descriptive tags to their favorite Web sites. In recent years, tagging has been adopted by many other social media sites to enable users to tag blogs (Technorati), images (Flickr), music (Last.fm), scientific papers (CiteULike), videos (YouTube), etc.

The distinguishing feature of tagging systems is that they use an uncontrolled vocabulary. This is in marked contrast to previous attempts to organize information via formal taxonomies and classification systems. A formal classification system, e.g., Linnaean classification of living things, puts an object in a unique place within a hierarchy. Thus, a **tiger** (*Panthera tigris*) is a carnivorous mammal that belongs to the genus *Panthera*, which also includes large cats, such as lions and leopards. Tiger is also part of the *felidae* family, which includes small cats, such as the familiar house cat of the genus *Felis*.

Tagging is a non-hierarchical and non-exclusive categorization, meaning that a user can choose to highlight any one of the tagged object's facets or properties. Adapting the example from Golder and Huberman (Golder and Huberman 2005), suppose a user takes an image of a Siberian tiger. Most likely, the user is not familiar with the formal name of the species (*P. tigris altaica*) and will tag it with the keyword "tiger." Depending on his needs or mood, the user may even tag it with more general or specific terms, such as "animal," "mammal" or "Siberian." The user may also note that the image was taken at the "zoo" and that he used his "telephoto" lens to get the shot. Rather than forcing the image into a hierarchy or multiple hierarchies based on the equipment used to take the photo, the place where the image was taken, type of animal depicted, or even the animal's provenance, tagging system allows the user to locate the image by any of its properties by filtering the entire image set on any of the tags. Thus, searching on the tag "tiger" will return all the images of tigers the user has taken, including Siberian and Bengal tigers, while searching on "Siberian" will return the images of Siberian animals, people or artifacts the user has photographed. Filtering on both "Siberian" and "tiger" tags will return the intersection of the images tagged with those keywords, in other words, the images of Siberian tigers.

As Golder and Huberman point out, tagging systems are vulnerable to problems that arise when users try to attach semantics to objects through keywords. These problems are exacerbated in social media where users may use different tagging conventions, but still want to take advantage of the others' tagging activities. The first problem is of homonymy, where the same tag may have different meanings. For example, the "tiger" tag could be applied to the mammal or to Apple computer's operating system. Searching on the tag "tiger" will return many images unrelated the carnivorous mammals, requiring the user to sift through possibly a large amount of irrelevant content. Another problem related to homonymy is that of polysemy, which arises when a word has multiple related meanings, such as "apple" to mean the company or any of its products. Another problem is that

of synonymy, or multiple words having the same or related meaning, for example, "baby" and "infant." The problem here is that if the user wants all images of young children in their first year of life, searching on the tag "baby" may not return all relevant images, since other users may have tagged similar photographs with "infant." Of course, plurals ("tigers" vs "tiger") and many other tagging idiosyncrasies ("myson" vs "son") may also confound a tagging system.

Golder and Huberman identify yet another problem that arises when using tags for categorization — that of the "basic level." A given item can be described by terms along a spectrum of specificity, ranging from specific to general. A Siberian tiger can be described as a "tiger," but also as a "mammal" and "animal." The basic level is the category people choose for an object when communicating to others about it. Thus, for most people, the basic level for canines is "dog," not the more general "animal" or the more specific "beagle." However, what constitutes the basic level varies between individuals, and to a large extent depends on the degree of expertise. To a dog expert, the basic level may be the more specific "beagle" or "poodle," rather than "dog." The basic level problem arises when different users choose to describe the item at different levels of specificity. For example, a dog expert tags an image of a beagle as "beagle," whereas the average user may tag a similar image as "dog." Unless the user is aware of the basic level variation and supplies more specific (and more general) keywords during tag search, he may miss a large number of relevant images.

Despite these problems, tagging is a light weight, flexible categorization system. The growing number of tagged images provides evidence that users are adopting tagging on Flickr (Marlow *et al.* 2006). There is speculation (Mika 2005) that collective tagging will lead to a common informal classification system, dubbed a "folksonomy," that will be used to organize all information from all users. Developing value-added systems on top of tags, e.g., which allow users to better browse or search for relevant items, will only accelerate wider acceptance of tagging.

Anatomy of Flickr

Flickr consists of a collection of interlinked user, photo, tag and group pages. A typical Flickr photo page is shown in Figure 1. It provides a variety of information about the image: who uploaded it and when, what groups it has been submitted to, its tags, who commented on the image and when, how many times the image was viewed or bookmarked as a "favorite." Clicking on a user's name brings up that user's photo stream, which shows the latest photos she has uploaded, the images she marked as "favorite," and her profile, which gives information about the user, including a list of her contacts and groups she belong to. Clicking on the tag shows user's images that have been tagged with this keyword, or all public images that have been similarly tagged. Finally, the group link brings up the group's page, which shows the photo group, group membership, popular tags, discussions and other information about the group.

¹<http://del.icio.us>

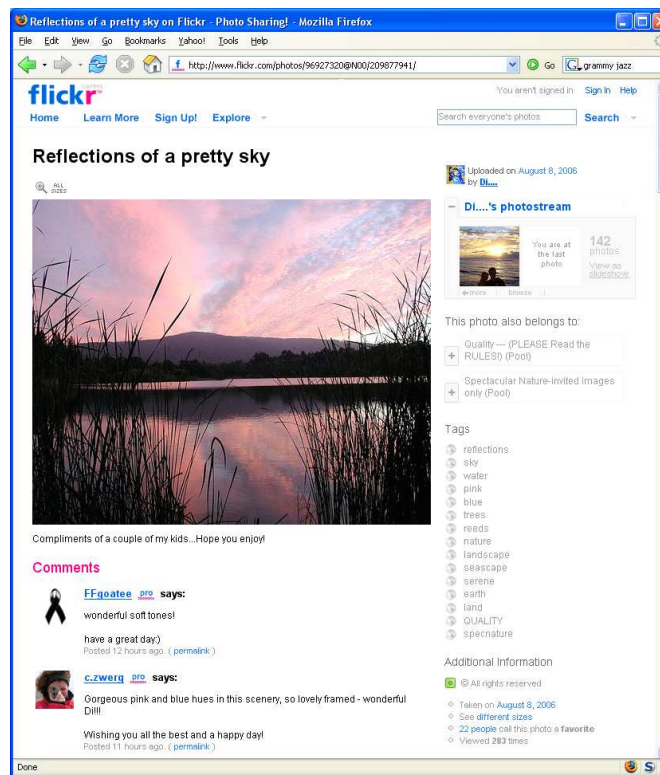


Figure 1: A typical photo page on Flickr

Groups Flickr allows users to create special interest groups on any imaginable topic. There are groups for showcasing exceptional images, group for images of circles within a square, groups for closeups of flowers, for the color red (and every other color and shade), groups for rating submitted images, or those used solely to generate comments. Some groups are even set up as games, such as *The Infinite Flickr*, where the rule is that a user post an image of herself looking at the screen showing the last image (of a user looking at the screen showing next to last image, etc).

There is redundancy and duplication in groups. For example, groups for child photography include *Children's Portraits*, *Kidpix*, *Flickr's Cutest Kids*, *Kids in Action*, *Toddlers*, etc. A user chooses one, or usually several, groups to which to submit an image. We believe that group names can be viewed as a kind of publicly agreed upon tags.

Contacts Flickr allows users to designate others as friends or contacts and makes it easy to track their activities. A single click on the "Contacts" hyperlink shows the user the latest images from his or her contacts. Tracking activities of friends is a common feature of many social media sites and is one of their major draws.

Interestingness Flickr uses the "interestingness" criterion to evaluate the quality of the image. Although the algorithm that is used to compute this is kept secret to prevent gaming the system, certain metrics are taken into account: "where

the clickthroughs are coming from; who comments on it and when; who marks it as a favorite; its tags and many more things which are constantly changing."²

Browsing and searching

Flickr offers the user a number of browsing and searching methods. One can browse by popular tags, through the groups directory, through the Explore page and the calendar interface, which provides access to the 500 most "interesting" images on any given day. A user can also browse geotagged images through the recently introduced map interface. Finally, Flickr allows for social browsing through the "Contacts" interface that shows in one place the recent images uploaded by the user's designated contacts.

Flickr allows searching for photos using full text or tag search. A user can restrict the search to all public photos, his or her own photos, photos she marked as her favorite, or photos from a specific contact. The advanced search interface currently allows further filtering by content type, date and camera.

Search results are by default displayed in reverse chronological order of being uploaded, with the most recent images on top. Another available option is to display images by their "interestingness" value, with the most "interesting" images on top.

²<http://flickr.com/explore/interesting/>

Personalizing search results

Suppose a user is interested in wildlife photography and wants to see images of tigers on Flickr. The user can search for all public images tagged with the keyword “tiger.” As of March 2007, such a search returns over 55,500 results. When images are arranged by their “interestingness,” the first page of results contains many images of tigers, but also of a tiger shark, cats, butterfly and a fish. Subsequent pages of search results show, in addition to tigers, children in striped suits, flowers (tiger lily), more cats, Mac OS X (tiger) screenshots, golfing pictures (Tiger Woods), etc. In other words, results include many false positives, images that are irrelevant to what the user had in mind when executing the search.

We assume that when the search term is ambiguous, the sense that the user has in mind is related to her interests. For example, when a child photographer is searching for pictures of a “newborn,” she is most likely interested in photographs of human babies, not kittens, puppies, or ducklings. Similarly, a nature photographer specializing in macro photography is likely to be interested in insects when searching on the keyword “beetle,” not a Volkswagen car. Users express their photography preferences and interests in a number of ways on Flickr. They express them through their contacts (photographers they choose to watch), through the images they upload to Flickr, through the tags they add to these images, through the groups they join, and through the images of other photographers they mark as their favorite. In this paper we show that we can personalize results of tag search by exploiting information about user’s preferences. In the sections below, we describe two search personalization methods: one that relies on user-created tags and one that exploits user’s contacts. We show that both methods improve search performance by reducing the number of false positives, or irrelevant results, returned to the user.

Data collections

To show how user-created metadata can be used to personalize results of tag search, we retrieved a variety of data from Flickr using their public API.

Data sets

We collected images by performing a single keyword tag search of all public images on Flickr. We specified that the returned images are ordered by their “interestingness” value, with most interesting images first. We retrieved the links to the top 4500 images for each of the following search terms:

tiger possible senses include (a) big cat (e.g., Asian tiger), (b) shark (Tiger shark), (c) flower (Tiger Lily), (d) golfing (Tiger Woods), etc.

newborn possible senses include (a) a human baby, (b) kitten, (c) puppy, (d) duckling, (e) foal, etc.

beetle possible senses include (a) a type of insect and (b) Volkswagen car model

For each image in the set, we used Flickr’s API to retrieve the name of the user who posted the image (image owner), and all the image’s tags and groups.

| query | relevant | not relevant | precision |
|---------|----------|--------------|-----------|
| newborn | 412 | 83 | 0.82 |
| tiger | 337 | 156 | 0.67 |
| beetle | 232 | 268 | 0.46 |

Table 1: Relevance results for the top 500 images retrieved by tag search

Users

Our objective is to personalize tag search results; therefore, to evaluate our approach, we need to have users to whose interests the search results are being tailored. We identified four users who are interested in the first sense of each search term. For the *newborn* data set, those users were one of the authors of the paper and three other contacts within that user’s social network who were known to be interested in child photography. For the other datasets, the users were chosen from among the photographers whose images were returned by the tag search. We studied each user’s profile to confirm that the user was interested in that sense of the search term. We specifically looked at group membership and user’s tags. Thus, for the *tiger* data set, groups that pointed to the user’s interest in *P. tigris* were **Big Cats**, **Zoo**, **The Wildlife Photography**, etc. In addition to group membership, tags that pointed to user’s interest in a topic, e.g., for the *beetle* data set, we assumed that users who used tags *nature* and *macro* were interested in insects rather than cars. Likewise, for the *newborn* data set, users who had uploaded images they tagged with *baby* and *child* were probably interested in human newborns.

For each of the twelve users, we collected the names of their contacts, or **Level 1 contacts**. For each of these contacts, we also retrieved the list of their contacts. These are called **Level 2 contacts**. In addition to contacts, we also retrieved the list of all the tags, and their frequencies, that the users had used to annotate their images. In addition to **all tags**, we also extracted a list of **related tags** for each user. These are the tags that appear together with the tag used as the search term in the user’s photos. In other words, suppose a user, who is a child photographer, had used tags such as “baby”, “child”, “newborn”, and “portrait” in her own images. Tags related to *newborn* are all the tags that co-occur with the “newborn” tag in the user’s own images. This information was also extracted via Flickr’s API.

Search results

We manually evaluated the top 500 images in each data set and marked each as relevant if it was related to the first sense of the search term listed above, not relevant or undecided, if the evaluator could not understand the image well enough to judge its relevance.

In Table 1, we report the precision of the search within the 500 labeled images, as judged from the point of view of the searching users. Precision is defined as the ratio of relevant images within the result set over the 500 retrieved images. Precision of tag search on these sample queries is not very high due to the presence of false positives — images not relevant to the sense of the search term the user had in mind.

In the sections below we show how to improve search performance by taking into consideration supplementary information about user’s interests provided by her contacts and tags.

Personalizing by contacts

Flickr encourages users to designate others as contacts by making it easy to view the latest images submitted by them through the “Contacts” interface. Users add contacts for a variety of reasons, including keeping in touch with friends and family, as well as to track photographers whose work is of interest to them. We claim that the latter reason is the most dominant of the reasons. Therefore, we view user’s contacts as an expression of the user’s interests. In this section we show that we can improve tag search results by filtering through the user’s contacts. To personalize search results for a particular user, we simply restrict the images returned by the tag search to those created by the user’s contacts.

Table 2 shows how many of the 500 images in each data set came from a user’s contacts. The column labeled “# L1” gives the number of user’s Level 1 contacts. The following columns show how many of the images were marked as relevant or not relevant by the filtering method, as well as precision and recall relative to the 500 images in each data set. Recall measures the fraction of relevant retrieved images relative to all relevant images within the data set. The last column “improv” shows percent improvement in precision over the plain (unfiltered) tag search.

As Table 2 shows, filtering by contacts improves the precision of tag search for most users anywhere from 22% to over 100% when compared to plain search results in Table 1. The best performance is attained for users within the *newborn* set, with a large number of relevant images correctly identified as being relevant, and no irrelevant images admitted into the result set. The *tiger* set shows an average precision gain of 42% over four users, while the *beetle* set shows an 85% gain.

Increase in precision is achieved by reducing the number of false positives, or irrelevant images that are marked as relevant by the search method. Unfortunately, this gain comes at the expense of recall: many relevant images are missed by this filtering method. In order to increase recall, we enlarge the contacts set by considering two levels of contacts: user’s contacts (Level 1) and her contacts’ contacts (Level 2). The motivation for this is that if the contact relationship expresses common interests among users, user’s interests will also be similar to those of her contacts’ contacts.

The second half of Table 2 shows the performance of filtering the search results by the combined set of user’s Level 1 and Level 2 contacts. This method identifies many more relevant images, although it also admits more irrelevant images, thereby decreasing precision. This method still shows precision improvement over plain search, with precision gain of 9%, 16% and 11% respectively for the three data sets.

Personalizing by tags

In addition to creating lists of contacts, users express their photography interests through the images they post on Flickr. We cannot yet automatically understand the content of images. Instead, we turn to the metadata added by the user to the image to provide a description of the image. The metadata comes in a variety of forms: image title, description, comments left by other users, tags the image owner added to it, as well as the groups to which she submitted the image. As we described in the paper, tags are useful image descriptors, since they are used to categorize the image. Similarly, group names can be viewed as public tags that a community of users have agreed on. Submitting an image to a group is, therefore, equivalent to tagging it with a public tag.

In the section below we describe a probabilistic model that takes advantage of the images’ tag and group information to discover latent topics in each search set. The users’ interests can similarly be described by collections of tags they had used to annotate their own images. The latent topics found by the model can be used to personalize search results by finding images on topics that are of interest to a particular user.

Model definition

We need to consider four types of entities in the model: a set of users $U = \{u_1, \dots, u_n\}$, a set of images or photos $I = \{i_1, \dots, i_m\}$, a set of tags $T = \{t_1, \dots, t_o\}$, and a set of groups $G = \{g_1, \dots, g_p\}$. A photo i_x posted by owner u_x is described by a set of tags $\{t_{x1}, t_{x2}, \dots\}$ and submitted to several groups $\{g_{x1}, g_{x2}, \dots\}$. The post could be viewed as a tuple $\langle i_x, u_x, \{t_{x1}, t_{x2}, \dots\}, \{g_{x1}, g_{x2}, \dots\} \rangle$. We assume that there are n users, m posted photos and p groups in Flickr. Meanwhile, the vocabulary size of tags is q . In order to filter images retrieved by Flickr in response to tag search and personalize them for a user u , we compute the conditional probability $p(i|u)$, that describes the probability that the photo i is relevant to u based on her interests. Images with high enough $p(i|u)$ are then presented to the user as relevant images.

As mentioned earlier, users choose tags from an uncontrolled vocabulary according to their styles and interests. Images of the same subject could be tagged with different keywords although they have similar meaning. Meanwhile, the same keyword could be used to tag images of different subjects. In addition, a particular tag frequently used by one user may have a different meaning to another user. Probabilistic models offer a mechanism for addressing the issues of synonymy, polysemy and tag sparseness that arise in tagging systems.

We use a probabilistic topic model (Rosen-Zvi *et al.* 2004) to model user’s image posting behavior. As in a typical probabilistic topic model, topics are hidden variables, representing knowledge categories. In our case, topics are equivalent to image owner’s interests. The process of photo posting by a particular user could be described as a stochastic process:

- User u decides to post a photo i .

| user | # L1 | rel. | not rel. | Pr | Re | improv | # L2+L2 | rel. | not rel. | Pr | Re | improv |
|----------------|------|------|----------|------|------|--------|---------|------|----------|------|------|--------|
| newborn | | | | | | | | | | | | |
| user1 | 719 | 232 | 0 | 1.00 | 0.56 | 22% | 49,539 | 349 | 62 | 0.85 | 0.85 | 4% |
| user2 | 154 | 169 | 0 | 1.00 | 0.41 | 22% | 10,970 | 317 | 37 | 0.9 | 0.77 | 10% |
| user3 | 174 | 147 | 0 | 1.00 | 0.36 | 22% | 13,153 | 327 | 39 | 0.89 | 0.79 | 9% |
| user4 | 128 | 132 | 0 | 1.00 | 0.32 | 22% | 8,439 | 310 | 29 | 0.91 | 0.75 | 11% |
| tiger | | | | | | | | | | | | |
| user5 | 63 | 11 | 1 | 0.92 | 0.03 | 37% | 13,142 | 255 | 71 | 0.78 | 0.76 | 16% |
| user6 | 103 | 78 | 3 | 0.96 | 0.23 | 44% | 14,425 | 266 | 83 | 0.76 | 0.79 | 13% |
| user7 | 62 | 65 | 1 | 0.98 | 0.19 | 47% | 7,270 | 226 | 60 | 0.79 | 0.67 | 18% |
| user8 | 56 | 30 | 0 | 0.97 | 0.09 | 44% | 7,073 | 240 | 63 | 0.79 | 0.71 | 18% |
| beetle | | | | | | | | | | | | |
| user9 | 445 | 18 | 1 | 0.95 | 0.08 | 106% | 53,480 | 215 | 221 | 0.49 | 0.93 | 7% |
| user10 | 364 | 35 | 8 | 0.81 | 0.15 | 77% | 41,568 | 208 | 217 | 0.49 | 0.90 | 7% |
| user11 | 783 | 78 | 25 | 0.75 | 0.34 | 65% | 62,610 | 218 | 227 | 0.49 | 0.94 | 7% |
| user12 | 102 | 7 | 1 | 0.88 | 0.03 | 90% | 14,324 | 163 | 152 | 0.52 | 0.70 | 13% |

Table 2: Results of filtering tag search by user’s contacts. “# L1” denotes the number of Level 1 contacts and “# L1+L2” shows the number of Level 1 and Level 2 contacts, with the succeeding columns displaying filtering results of that method: the number of images marked relevant or not relevant, as well as precision and recall of the filtering method relative to the top 500 images. The columns marked “improv” show improvement in precision over plain tag search results.




Figure 2: Graphical representation for model-based information filtering. U , T , G and Z denote variables “User”, “Tag”, “Group”, and “Topic” respectively. N_t represents a number of tag occurrences for a one photo (by the photo owner); D represents a number of all photos on Flickr. Meanwhile, N_g denotes a number of groups for a particular photo.

- Based on user u ’s interests and the subject of the photo, a set of topics z are chosen.
- Tag t is then selected based on the set of topics chosen in the previous state.
- In case that u decides to expose her photo to some groups, a group g is then selected according to the chosen topics.

The process is depicted in a graphical form in Figure 2. We do not treat the image i as a variable in the model but view it as a co-occurrence of a user, a set of tags and a set of groups. From the process described above, we can represent the joint probability of user, tag and group for a particular photo as

$$\begin{aligned}
 p(i) &= p(u_i, T_i, G_i) \\
 &= p(u_i) \cdot \left(\prod_{n_t} \left(\sum_k p(z_k | u_i) p(t_i | z) \right)^{n_i(t)} \right) \\
 &\quad \cdot \left(\prod_{n_g} \left(\sum_k p(z_k | u_i) p(g_i | z) \right)^{n_i(g)} \right).
 \end{aligned}$$

Note that it is straightforward to exclude photo’s group information from the above equation simply by omitting the terms relevant to g . n_t and n_g is a number of all possible tags and groups respectively in the data set. Meanwhile, $n_i(t)$ and $n_i(g)$ act as indicator functions: $n_i(t) = 1$ if an image i is tagged with tag t ; otherwise, it is 0. Similarly, $n_i(g) = 1$ if an image i is submitted to group g ; otherwise, it is 0. k is the predefined number of topics.

The joint probability of photos in the data set I is defined as

$$p(I) = \prod_m p(i_m).$$

In order to estimate parameters $p(z|u_i)$, $p(t_i|z)$, and $p(g_i|z)$, we define a log likelihood L , which measures how the estimated parameters fit the observed data. According to the EM algorithm (Dempster *et al.* 1977), L will be used as an objective function to estimate all parameters. L is defined as

$$L(I) = \log(p(I)).$$

In the expectation step (E-step), the joint probability of the hidden variable Z given all observations is computed from the following equations:

$$p(z|t, u) \propto p(z|u) \cdot p(t|z) \quad (1)$$

$$p(z|g, u) \propto p(z|u) \cdot p(g|z). \quad (2)$$

L cannot be maximized easily, since the summation over the hidden variable Z appears inside the logarithm. We instead maximize the expected complete data log-likelihood over the hidden variable, $E[L^c]$, which is defined as

$$\begin{aligned} E[L^c] &= \sum_i \log(p(u_i)) \\ &+ \sum_i \sum_t n_i(t) \cdot \sum_z p(z|u, t) (\log(p(z|u) \cdot (t|z))) \\ &+ \sum_i \sum_g n_i(g) \cdot \sum_z p(z|u, g) (\log(p(z|g) \cdot (g|z))) \end{aligned}$$

Since the term $\sum_i \log(p(u_i))$ is not related to parameters and can be computed directly from the observed data, we discard this term from the expected complete data log-likelihood. With normalization constraints on all parameters, Lagrange multipliers τ, ρ, ψ are added to the expected log likelihood, yielding the following equation

$$\begin{aligned} H = E[L^c] &+ \sum_z \tau_z \left(1 - \sum_t p(t|z)\right) \\ &+ \sum_z \rho_z \left(1 - \sum_g p(g|z)\right) \\ &+ \sum_u \psi_u \left(1 - \sum_z p(z|u)\right). \end{aligned}$$

We maximize H with respect to $p(t|z_k)$, $p(g|z_k)$, and $p(z_k|u)$, and then eliminate the Lagrange multipliers to obtain the following equations for the maximization step:

$$p(t|z) \propto \sum_m n_i(t) \cdot p(z|t, u) \quad (3)$$

$$p(g|z) \propto \sum_m n_i(g) \cdot p(z|g, u) \quad (4)$$

$$\begin{aligned} p(z_k|u_m) &\propto \sum_m \left(\sum_t n_m(t) \cdot p(z_k|u_m, t) \right. \\ &\quad \left. + \sum_g n_m(g) \cdot p(z_k|u_m, g) \right). \end{aligned} \quad (5)$$

The algorithm iterates between E and M step until the log likelihood for all parameter values converge.

Model-based personalization

We can use the model developed in the previous section to find the images i most relevant to the interests of a particular user u' . We do so by learning the parameters of the model from the data and using these parameters to compute the conditional probability $p(i|u')$. This probability can be factorized as follows:

$$p(i|u') = \sum_z p(u_i, T_i, G_i|z) \cdot p(z|u'), \quad (6)$$

where u_i is the owner of image i in the data set, and T_i and G_i are, respectively, the set of all the tags and groups for the image i .

The former term in Equation 6 can be factorized further as

$$\begin{aligned} p(u_i, T_i, G_i|z) &\propto p(T_i|z) \cdot (G_i|z) \cdot (z|u_i) \cdot p(u_i) \\ &= \left(\prod_{t_i} p(t_i|z)\right) \cdot \left(\prod_{g_i} p(g_i|z)\right) \cdot p(z|u_i) \cdot p(u_i). \end{aligned}$$

We can use the learned parameters to compute this term directly.

We represent the interests of user u' as an aggregate of the tags that u' had used in the past for tagging her own images. This information is used to approximate $p(z|u')$:

$$p(z|u') \propto \sum_t n(t' = t) \cdot p(z|t)$$

where $n(t' = t)$ is a frequency (or weight) of tag t' used by u' . Here we view $n(t' = t)$ is proportional to $p(t'|u')$. Note that we can use either all the tags u' had applied to the images in her photostream, or a subset of these tags, e.g., only those that co-occur with some tag in user's images.

Evaluation

We trained the model separately on each data set of 4500 images. We fixed the number of topics at ten. We then evaluated our model-based personalization framework by using the learned parameters and the information about the interests of the selected users to compute $p(i|u')$ for the top 500 (manually labeled) images in the set. Information about user's interests was captured either by (1) *all tags* (and their frequencies) that are used in all the images of the user's photostream or (2) *related tags* that occurred in images that were tagged with the search keyword (e.g., "newborn") by the user.

Computation of $p(t|z)$ is central to the parameter estimation process, and it tells us something about how strongly a tag t contributes to a topic z . Table 3 shows the most probable 25 tags for each topic for the *tiger* data set trained on ten topics. Although the tag "tiger" dominates most topics, we can discern different themes from the other tags that appear in each topic. Thus, topic z_5 is obviously about domestic cats, while topic z_8 is about Apple computer products. Meanwhile, topic z_2 is about flowers and colors ("flower," "lily," "yellow," "pink," "red"); topic z_6 is about about places ("losangeles," "sandiego," "lasvegas," "stuttgart,"), presumably because these places have zoos. Topic z_7 contains several variations of tiger's scientific name, "panthera tigris." This method appears to identify related words well. Topic z_5 , for example, gives synonyms "cat," "kitty," as well as the more general term "pet" and the more specific terms "kitten" and "tabby." It even contains the Spanish version of the word: "gatto." In future work we plan to explore using this method to categorize photos in a more abstract way. We also note that related terms can be used to increase search recall by providing additional keywords for queries.

Table 4 presents results of model-based personalization for the case that uses information from all of user's tags. The model was trained with ten topics. Results are presented for different thresholds. The first two columns, for example, report precision and recall for a high threshold that

| z₁ | z₂ | z₃ | z₄ | z₅ |
|--|---|--|---|--|
| tiger zoo animal nature animals wild tijger wildlife ilovenature cub siberiantiger blijdorp london australia portfolio white dierentuin toronto stripes amurtiger nikonstunninggallery s5600 eyes sydney cat | tiger specanimal animalkingdomelite abigfave flower butterfly macro yellow swallowtail lily green canon insect nature pink red flowers orange eastern usa impressedbeauty tag2 specnature black streetart | tiger cat kitty cute kitten cats orange eyes pet tabby stripes whiskers white art feline fur animal gatto pets black paws furry nose teeth beautiful | tiger thailand bengal animals tigers canon d50 tigertemple 20d white nikon kanchanaburi detroit life michigan detroitzoo eos temple park asia ball marineworld baseball detroittigres wild | tiger cat animal animals zoo bigcat cats tigre animalplanet tigers bigcats whitetiger mammal wildlife colorado stripes denver sumatrantiger white feline mammals sumatran exoticcats exoticcat big |
| z₆ | z₇ | z₈ | z₉ | z₁₀ |
| tiger tigers dczoo tigercub california lion cat cc100 florida girl wilhelma self lasvegas stuttgart me baby tattoo endangered illustration ?? losangeles portrait sandiego lazoo giraffe | nationalzoo tiger sumatrantiger zoo nikon washingtondc smithsonian washington animals cat bigcat tigris panthera bigcats d70s pantheratigrissumatrae dc sumatrae animal 2005 pantheratigris nikond70 d70 2006 topv111 | tiger apple mac osx macintosh screenshot macosx desktop imac stevejobs dashboard macbook powerbook os 104 canon x ipod computer ibook intel keyboard widget wallpaper laptop | tiger india canon wildlife impressedbeauty endangered safari wildanimals wild tag1 tag3 park taggedout katze nature bravo nikon asia canonrebelxt bandhavgarh vienna schnbrunn zebra pantheratigris d2x | tiger lion dog shark nyc cat man people arizona rock beach sand sleeping tree forest puppy bird portrait marwell boy fish panther teeth brooklyn bahamas |

Table 3: Top tags ordered by $p(t=z)$ for the ten topic model of the “tiger” data set.

| | Pr | Re | Pr | Re | Pr | Re | Pr | Re | Pr | Re |
|----------------|------|------|-------|------|-------|------|--------|------|--------|------|
| newborn | | | | | | | | | | |
| | n=50 | | n=100 | | n=200 | | n=300 | | n=412* | |
| user1 | 1.00 | 0.12 | 1.00 | 0.24 | 1.00 | 0.49 | 0.94 | 0.68 | 0.89 | 0.89 |
| user2 | 1.00 | 0.12 | 1.00 | 0.24 | 1.00 | 0.49 | 0.92 | 0.67 | 0.87 | 0.87 |
| user3 | 1.00 | 0.12 | 0.88 | 0.21 | 0.84 | 0.41 | 0.85 | 0.62 | 0.89 | 0.89 |
| user4 | 1.00 | 0.12 | 0.99 | 0.24 | 1.00 | 0.48 | 0.94 | 0.69 | 0.89 | 0.89 |
| tiger | | | | | | | | | | |
| | n=50 | | n=100 | | n=200 | | n=300 | | n=337* | |
| user5 | 0.94 | 0.14 | 0.90 | 0.27 | 0.82 | 0.48 | 0.80 | 0.71 | 0.79 | 0.79 |
| user6 | 0.76 | 0.11 | 0.80 | 0.24 | 0.79 | 0.47 | 0.77 | 0.69 | 0.77 | 0.77 |
| user7 | 0.94 | 0.14 | 0.90 | 0.27 | 0.82 | 0.48 | 0.80 | 0.71 | 0.79 | 0.79 |
| user8 | 0.90 | 0.13 | 0.88 | 0.26 | 0.82 | 0.49 | 0.79 | 0.71 | 0.79 | 0.79 |
| beetle | | | | | | | | | | |
| | n=50 | | n=100 | | n=200 | | n=232* | | n=300 | |
| user9 | 1.00 | 0.22 | 0.99 | 0.43 | 0.77 | 0.66 | 0.70 | 0.70 | 0.66 | 0.85 |
| user10 | 0.98 | 0.21 | 0.99 | 0.43 | 0.77 | 0.66 | 0.70 | 0.70 | 0.66 | 0.85 |
| user11 | 0.98 | 0.21 | 0.93 | 0.40 | 0.50 | 0.43 | 0.51 | 0.51 | 0.50 | 0.65 |
| user12 | 1.00 | 0.22 | 0.99 | 0.43 | 0.77 | 0.66 | 0.70 | 0.70 | 0.66 | 0.85 |

Table 4: Filtering results where a number of learned topics is 10, excluding group information, and user’s personal information obtained from all tags she used for her photos. Asterisk denotes R-precision of the method, or precision of the first n results, where n is the number of relevant results in the data set.

marks only the 50 most probable images as relevant. The remaining 450 images are marked as not relevant to the user. Recall is low, because many relevant images are excluded from the results for such a high threshold. As the threshold is decreased ($n = 100, n = 200, \dots$), recall relative to the 500 labeled images increases. Precision remains high in all cases, and higher than precision of the plain tag search reported in Table 1. In fact, most of the images in the top 100 results presented to the user are relevant to her query. The column marked with the asterisk gives the R-precision of the method, or precision of the first R results, where R is the number of relevant results. The average R-precision of this filtering method is 8%, 17% and 42% better than plain search precision on our three data sets.

Performance results of the approach that uses related tags instead of all tags are given in Table 5. We explored this direction, because we believed it could help discriminate between different topics that interest a user. Suppose, a child photographer is interested in nature photography as well as child portraiture. The subset of tags he used for tagging his “newborn” portraits will be different from the tags used for tagging nature images. These tags could be used to differentiate between newborn baby and newborn colt images. However, on the set of users selected for our study, using related tags did not appear to improve results. This could be because the tags a particular user used together with, for example, “beetle” do not overlap significantly with the rest of the data set.

Including group information did not significantly improve results (not presented in this manuscript). In fact, group information sometimes hurts the estimation rather than helps. We believe that this is because our data sets (sorted by Flickr according to image interestingness) are biased by the presence of general topic groups (e.g., **Search the Best, Spec-**

tacular Nature, Let’s Play Tag, etc.). We postulate that group information would help estimate $p(i|z)$ in cases where the photo has few or no tags. Group information would help filling in the missing data by using group name as another tag. We also trained the model on the data with 15 topics, but found no significant difference in results.

Previous research

Recommendation or personalization systems can be categorized into two main categories. One is collaborative filtering (Breese *et al.* 1998) which exploits item ratings from many users to recommend items to other like-minded users. The other is content-based recommendation, which relies on the contents of an item and user’s query, or other user information, for prediction (Mooney and Roy 2000). Our first approach, filtering by contacts, can be viewed as implicit collaborative filtering, where the user–contact relationship is viewed as a preference indicator: it assumes that the user likes all photos produced by her contacts. In our previous work, we showed that users do indeed agree with the recommendations made by contacts (Lerman 2007; Lerman and Jones 2007). This is similar to the ideas implemented by MovieTrust (Golbeck 2006), but unlike that system, social media sites do not require users to rate their trust in the contact.

Meanwhile, our second approach, filtering by tags (and groups), shares some characteristics with both methods. It is similar to collaborative filtering, since we use tags to represent agreement between users. It is also similar to content-based recommendation, because we represent image content by the tags and group names that have been assigned to it by the user.

Our model-based filtering system is technically similar to, but conceptually different from, probabilistic models pro-

| | Pr | Re | Pr | Re | Pr | Re | Pr | Re | Pr | Re |
|----------------|------|------|-------|------|-------|------|--------|------|--------|------|
| newborn | | | | | | | | | | |
| | n=50 | | n=100 | | n=200 | | n=300 | | n=412* | |
| | Pr | Re | Pr | Re | Pr | Re | Pr | Re | Pr | Re |
| user1 | 0.8 | 0.10 | 0.78 | 0.19 | 0.79 | 0.38 | 0.77 | 0.56 | 0.79 | 0.79 |
| user2 | 0.8 | 0.10 | 0.82 | 0.20 | 0.80 | 0.39 | 0.77 | 0.56 | 0.83 | 0.83 |
| user3 | 0.98 | 0.12 | 0.88 | 0.21 | 0.84 | 0.41 | 0.80 | 0.58 | 0.85 | 0.85 |
| user4 | 0.98 | 0.12 | 0.88 | 0.21 | 0.84 | 0.41 | 0.85 | 0.62 | 0.88 | 0.88 |
| tiger | | | | | | | | | | |
| | n=50 | | n=100 | | n=200 | | n=300 | | n=337* | |
| | Pr | Re | Pr | Re | Pr | Re | Pr | Re | Pr | Re |
| user5 | 0.84 | 0.12 | 0.86 | 0.26 | 0.78 | 0.46 | 0.78 | 0.69 | 0.77 | 0.77 |
| user6 | 0.72 | 0.11 | 0.79 | 0.23 | 0.78 | 0.46 | 0.76 | 0.68 | 0.76 | 0.76 |
| user7 | 0.72 | 0.11 | 0.78 | 0.23 | 0.78 | 0.46 | 0.76 | 0.68 | 0.76 | 0.76 |
| user8 | 0.9 | 0.13 | 0.82 | 0.24 | 0.80 | 0.47 | 0.78 | 0.69 | 0.78 | 0.78 |
| beetle | | | | | | | | | | |
| | n=50 | | n=100 | | n=200 | | n=232* | | n=300 | |
| | Pr | Re | Pr | Re | Pr | Re | Pr | Re | Pr | Re |
| user9 | 0.78 | 0.17 | 0.62 | 0.27 | 0.58 | 0.50 | 0.54 | 0.54 | 0.53 | 0.68 |
| user10 | 0.98 | 0.21 | 0.88 | 0.38 | 0.77 | 0.66 | 0.72 | 0.72 | 0.65 | 0.84 |
| user11 | 0.96 | 0.21 | 0.74 | 0.32 | 0.62 | 0.53 | 0.59 | 0.59 | 0.56 | 0.72 |
| user12 | 0.98 | 0.21 | 0.99 | 0.43 | 0.77 | 0.66 | 0.70 | 0.70 | 0.66 | 0.85 |

Table 5: Filtering results where a number of learned topics is 10, excluding group information, and user’s personal information obtained from all tags she used for her photos, which are tagged by the search term

posed by (Popescul *et al.* 2001). Both models are probabilistic generative models that describe co-occurrences of users and items of interest. In particular, the model assumes a user generates her topics of interest; then the topics generate documents and words in those documents if the user prefers those documents. In our model, we metaphorically assume the photo owner generates her topics of interest. The topics, in turn, generate tags that the owner used to annotate her photo. However, unlike the previous work, we do not treat photos as variables, as they do for documents. This is because images are tagged only by their owners; meanwhile, in their model, all users who are interested in a document generate topics for that document.

Our model-based approach is almost identical to the author-topic model (Rosen-Zvi *et al.* 2004). However, we extend their framework to address (1) how to exploit photo’s group information for personalized information filtering; (2) how to approximate user’s topics of interest from partially observed personal information (the tags the user used to describe her own images). For simplicity, we use the classical EM algorithm to train the model; meanwhile they use a stochastic approximation approach due to the difficulty involved in performing exact an inference for their generative model.

Conclusions and future work

We presented two methods for personalizing results of image search on Flickr. Both methods rely on the meta-data users create through their everyday activities on Flickr, namely user’s contacts and the tags they used for annotating their images. We claim that this information captures user’s tastes and preferences in photography and can be used to personalize search results to the individual user. We showed

that both methods dramatically increase search precision. We believe that increasing precision is an important goal for personalization, because dealing with the information overload is the main issue facing users, and we can help users by reducing the number of irrelevant results the user has to examine (false positives). Having said that, our tag-based approach can also be used to expand the search by suggesting relevant related keywords (e.g., “pantheratigris,” “big-cat” and “cub” for the query *tiger*).

In addition to tags and contacts, there exists other meta-data, favorites and comments, that can be used to aid information personalization and discovery. In our future work we plan to address the challenge of combining these heterogeneous sources of evidence within a single approach. We will begin by combining contacts information with tags.

The probabilistic model needs to be explored further. Right now, there is no principled way to pick the number of latent topics that are contained in a data set. We also plan to have a better mechanism for dealing with uninformative tags and groups. We would like to automatically identify general interest groups, such as the *Let’s Play Tag* group, that do not help to discriminate between topics.

The approaches described here can be applied to other social media sites, such as Del.icio.us. We imagine that in near future, all of Web will be rich with metadata, of the sort described here, that will be used to personalize information search and discovery to the individual user.

Acknowledgements

This research is based on work supported in part by the National Science Foundation under Award Nos. IIS-0535182 and in part by DARPA under Contract No. NBCHD030010.

The U.S. Government is authorized to reproduce and dis-

tribute reports for Governmental purposes notwithstanding any copyright annotation thereon. The views and conclusions contained herein are those of the authors and should not be interpreted as necessarily representing the official policies or endorsements, either expressed or implied, of any of the above organizations or any person connected with them.

References

- [Breese *et al.* 1998] John Breese, David Heckerman, and Carl Kadie. Empirical analysis of predictive algorithms for collaborative filtering. In *Proceedings of the 14th Annual Conference on Uncertainty in Artificial Intelligence (UAI-98)*, pages 43–52, San Francisco, CA, 1998. Morgan Kaufmann.
- [Dempster *et al.* 1977] A. P. Dempster, N. M. Laird, and D. B. Rubin. Maximum likelihood from incomplete data via the em algorithm. *Journal of the Royal Statistical Society. Series B (Methodological)*, 39(1):1–38, 1977.
- [Golbeck 2006] J. Golbeck. Generating predictive movie recommendations from trust in social networks. In *Proceedings of the Fourth International Conference on Trust Management*, Pisa, Italy, May 2006.
- [Golder and Huberman 2005] S. A. Golder and B. A. Huberman. The structure of collaborative tagging systems. Technical report, HP Labs, 2005. <http://www.hpl.hp.com/research/idl/papers/tags/>.
- [Lerman and Jones 2007] K. Lerman and Laurie Jones. Social browsing on flickr. In *Proc. of International Conference on Weblogs and Social Media (ICWSM-07)*, 2007.
- [Lerman 2007] K. Lerman. Social networks and social information filtering on digg. In *Proc. of International Conference on Weblogs and Social Media (ICWSM-07)*, 2007.
- [Marlow *et al.* 2006] C. Marlow, M. Naaman, d. boyd, and M. Davis. Ht06, tagging paper, taxonomy, flickr, academic article, toread. In *Proceedings of Hypertext 2006*, New York, 2006. ACM, New York: ACM Press.
- [Mika 2005] P. Mika. Ontologies are us: A unified model of social networks and semantics. In *International Semantic Web Conference (ISWC-05)*, 2005.
- [Mooney and Roy 2000] Raymond J. Mooney and Loriene Roy. Content-based book recommending using learning for text categorization. In *Proceedings of 5th ACM Conference on Digital Libraries*, pages 195–204, San Antonio, US, 2000. ACM Press, New York, US.
- [Popescul *et al.* 2001] Alexandrin Popescul, Lyle Ungar, David Pennock, and Steve Lawrence. Probabilistic models for unified collaborative and content-based recommendation in sparse-data environments. In *17th Conference on Uncertainty in Artificial Intelligence*, pages 437–444, Seattle, Washington, August February–May 2001.
- [Rosen-Zvi *et al.* 2004] Michal Rosen-Zvi, Thomas Griffiths, Mark Steyvers, and Padhraic Smyth. The author-topic model for authors and documents. In *AUAI '04: Proceedings of the 20th conference on Uncertainty in artificial intelligence*, pages 487–494, Arlington, Virginia, United States, 2004. AUAI Press.